

Computer Systems II: Kernels and Memory Management

Noah Singer, George Klees

Montgomery Blair High School Computer Team

September 29, 2017

Overview

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

1 Kernels

2 Memory Management

Section 1: Kernels

Stored-program computers

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

- Programs compiled into a sequence of instructions called **machine code**
- Stored as simple data in RAM
 - Very long sequence of **bytes** (8 **bits**)
 - Each byte has a numerical **memory address**
 - CPU can **load and store** data from RAM
 - CPU also has much faster, smaller memory called **registers**
- Data executed as code by CPU

Operating systems

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

- Intermediary between hardware and applications
- Creates a high-level interface for application developers
- Controls access to hardware and enforces security procedures
- Often separated into core **kernel** and external **drivers**
 - Drivers are often used to interface with specific hardware or accomplish a specific task, and vary from computer to computer depending on hardware and setup
 - Kernel accomplishes core tasks regardless of specific hardware

Operating modes

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

- When the CPU first loads, any executed code can access all memory, CPU functions, and hardware devices
 - On a secure and reasonable system, applications must be **sandboxed**
- Modern CPUs have two modes that code can execute in
 - **Supervisor**: unrestricted access to resources, only granted to kernel which accomplishes core OS functions
 - **User**: access restricted to only certain memory, certain devices, etc. (for applications)
- Special instructions called **syscalls** allow a user mode application to “jump” into the kernel in supervisor mode to accomplish an OS task
 - Read a file from the hard drive
 - Create (fork) a new process

Types of kernels

- **Microkernel:** Only the bare minimum that requires supervisor mode is in the kernel
 - Many OS functions are user mode drivers
 - More secure and elegant
 - Slower, since many switches between user and supervisor modes may be required
- **Monolithic:** Most of the OS is in the kernel
 - Less secure since any vulnerability in the much larger kernel leads to supervisor control over the system
 - Faster and less complicated
- **Hybrid kernel:** Middle ground
 - Some user mode drivers, some supervisor mode drivers

Core kernel/OS functions

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

- **Memory management:** allocating and controlling memory for applications
- **Task management:** allowing multiple applications to run simultaneously
- Virtual filesystem (**VFS**): access to stored data
 - Many layers of drivers involved: storage drivers, bus drivers, filesystem drivers, etc.
 - Must present a uniform interface for applications
- Device and power management

Sample (simplified) x86 boot process

- 1 BIOS boot code is loaded from read-only memory (**ROM**), triggering code in hard drive Master Boot Record (**MBR**)
- 2 MBR triggers code in hard drive partition Volume Boot Record (**VBR**), loading the kernel and boot drivers from the filesystem with a filesystem driver
- 3 Kernel activates memory manager and enables paging
- 4 Kernel sets up syscalls and fault handlers
- 5 Kernel loads VFS using boot drivers and uses VFS to load other necessary drivers
- 6 Kernel initializes timers, various buses, and other hardware devices
- 7 Kernel begins multitasking (multicore?), switches into user mode, and loads user login application

Section 2: Memory Management

Memory management

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

- Every application's memory must be sandboxed from the others
- Each application has a unique **address space** of **virtual memory** mapped by the memory manager to **physical memory**
- When an application executes, it “sees” the address space of virtual memory
- In modern CPUs, memory is organized into medium-scale (often 4096 B) **pages**
- Task of kernel memory manager: for each application, create a correspondence between virtual pages and physical pages

Kernel memory manager

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

- **Physical memory allocator** allots pages of physical RAM to applications and the OS
- **Virtual memory manager** organizes the address space of applications
- Kernel communicates with paging hardware (memory management unit, or **MMU**) inside of CPU
 - **Page tables** stored in memory; MMU contains a pointer in a special register
 - MMU translates memory accesses from virtual to physical addresses
 - Lookups cached in the translation lookaside buffer (**TLB**)

Page tables

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

- Stored as an array of **page table entries (PTEs)**, each representing a page
- PTE stores detailed information, allowing for fine control over memory access:
 - Address of page in physical memory
 - Permissions
 - Caching information

Application address space

- Four main regions:
 - 1 Program code and data
 - 2 **Heap**: dynamically allocated data growing upwards
 - 3 Unallocated memory
 - 4 **Stack**: return addresses/stack frames, local variables, parameters, execution context
- When invalid memory is accessed, MMU throws a **page fault**, halting execution; can be used creatively to great advantage
 - Memory-mapped files
 - Swapping
- OS memory manager must: allocate and free memory, allow shared memory, map and unmap files, etc.

Conclusion

Computer
Systems II:
Kernels and
Memory
Management

Noah Singer,
George Klees

Kernels

Memory
Management

- OS interfaces between hardware and applications
- Code executes in user mode or supervisor mode
- Virtual memory allows kernel to control application address spaces