

Computer Systems III: Multitasking and Concurrency

Noah Singer, George Klees

Montgomery Blair High School Computer Team

October 5, 2017

Overview

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking

Concurrency

1 Multitasking

2 Concurrency

Section 1: Multitasking

Processes and threads

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking
Concurrency

- **Process:** A high-level task; an instance of a running program, with an address space and access to resources
- **Threads:** A low-level task; code running on the CPU, consisting of the actual execution state (CPU registers, stack, etc.)
- One or more threads per process
- **Multitasking:** CPU switches back and forth rapidly between threads

Program execution

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking

Concurrency

- 1 Create blank address space
- 2 Load program from disk into address space
- 3 Create initial thread with stack
- 4 Begin executing

Scheduling

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking

Concurrency

- **Scheduler** must select which threads to run and for how long
 - **Cooperative scheduling**: threads can run as long as they want until they yield
 - **Preemptive scheduling**: threads get a fixed amount of time (**quantum**)
- Threads are generally non-malicious; when any thread has no work to do, it often just yields to allow other threads a turn
- Threads often voluntarily yield when waiting for an operation to complete (**blocking**)
 - Most threads spend most of their time waiting for locks and I/O, so cooperative scheduling is common

Prioritization

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking

Concurrency

- Some threads are more important than others
- **High-priority** threads must complete immediately
 - Examples: GUI, some drivers
 - Often simply receive, route, and transmit requests, sleeping almost immediately after awaking
- **Low-priority** threads can be scheduled whenever there is no more important work
 - Examples: Background processes, virus scanners, updates

Task switching

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking

Concurrency

- When a thread exceeds its quantum, it is **preempted** and loses control
 - Triggered by a hardware timer
- Scheduler holds a queue of threads waiting to run
- When a yield or preemption occurs, the scheduler must switch **CPU contexts** (stack and registers)
 - 1 Save old CPU context
 - 2 Switch address spaces, etc. to new process
 - 3 Restore new CPU context

Section 2: Concurrency

Multiple CPUs

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking

Concurrency

- Scheduler also selects which CPU to schedule a thread on
 - CPU load, power management status, temperature, etc.
 - Due to caches and **affinity**, threads are better scheduled on CPUs on which they have already/recently executed
- **Concurrency** issues occur when threads on two CPUs access the same resource simultaneously
 - Or one thread is working, gets preempted, and then another thread accesses the same resource
 - The kernel is one such resource

Locks

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking

Concurrency

- **Spinlocks** simply record whether or not they are locked
 - Order in which threads attempt to acquire lock is irrelevant
 - Threads “spin” (in a while loop) while waiting for lock
- **Granularity**: the amount of code that is protected by locks
 - Too fine: doesn't fix the original issues
 - Too coarse: doesn't scale and difficult to develop/maintain

Atomic operations

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking
Concurrency

- One thread can be interrupted *while* acquiring the spinlock
- Uninterruptible operations; another CPU cannot interfere
- Thread cannot be preempted in the process of accessing a resource
- In multi-step atomic operations, memory bus is locked, slowing down the rest of the system
- **Compare and exchange** compares the value of `var` and `old`, setting `var` to `new` if it is equal to `old` and returning the original value of `var`

Issues with spinlocks

Computer
Systems III:
Multitasking
and
Concurrency

Noah Singer,
George Klees

Multitasking

Concurrency

- Threads cannot yield while waiting for spinlocks
- **Deadlock**: two threads both spinning to acquire resources the other thread has
- **Starvation**: a thread is continually outcompeted for locks, so it cannot access necessary resources
- Some more sophisticated locks solve these issues
 - **Ticket locks** á la deli counter or DMV